

Internet Protocols EBU5403

The Network Layer (Part II)

CI

Michael Chai (michael.chai@qmul.ac.uk)

Richard Clegg (r.clegg@qmul.ac.uk)

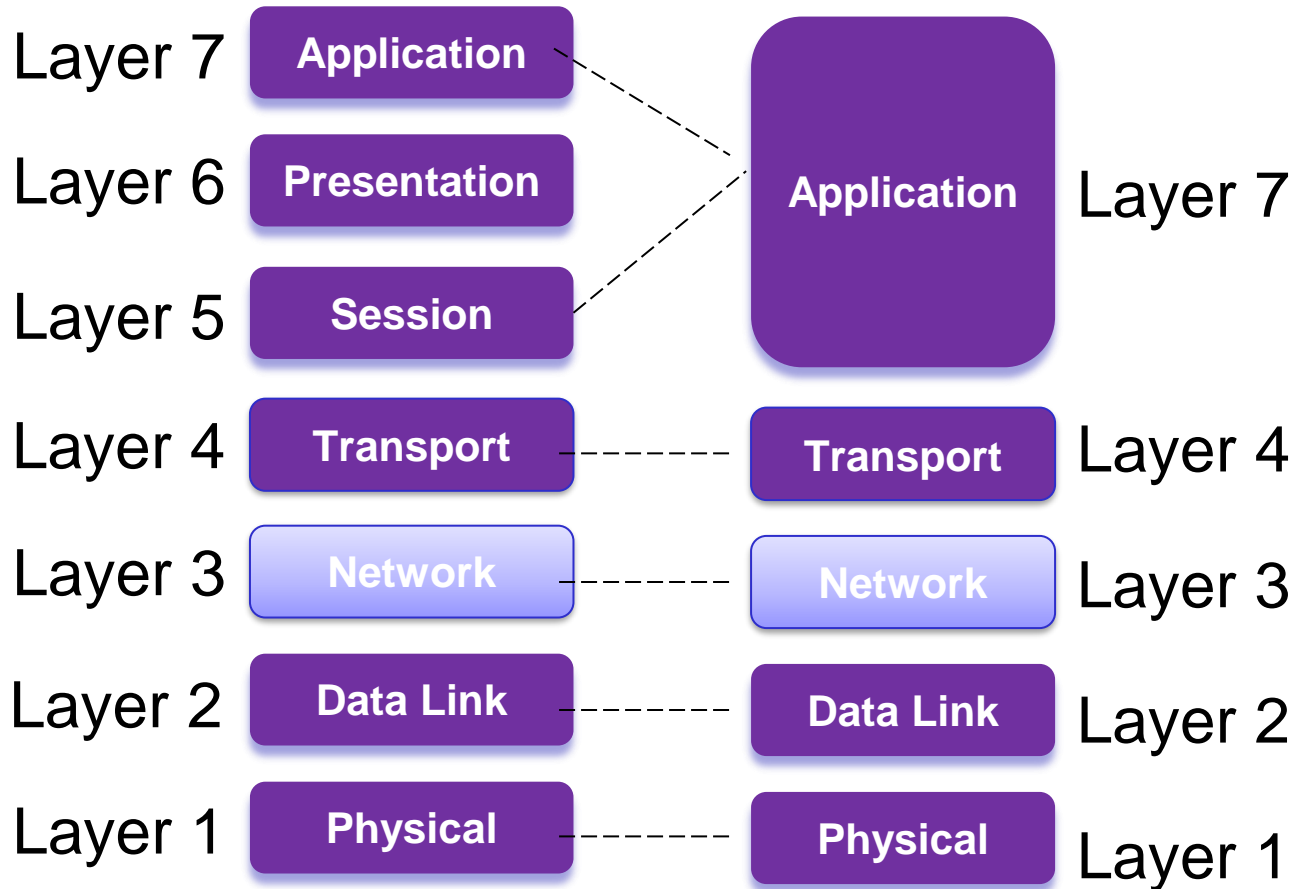
Cunhua Pan (c.pan@qmul.ac.uk)

| | Part 1 | Part 2 | Part 3 | Part 4 |
|------------------------|---------------|--------|------------|--------|
| Ecommerce + Telecoms 1 | Richard Clegg | | Cunhua Pan | |
| Telecoms 2 | Michael Chai | | | |

Structure of course

- Part 1
 - Introduction to IP Networks
 - The Transport layer (part I)
- Part 2
 - The Transport layer (part II)
 - The Network layer (part I)
- Part 3
 - The Network layer (part II)
 - The Data link layer (part I)
 - Router lab
- Part 4
 - The Data link layer (part II)
 - Network management and security
 - Class test

Network Layer



Network Data Plane: outline

4.1 Overview of Network layer

- data plane
- control plane

4.2 What's inside a router


4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing
- network address translation
- IPv6

4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

IPv4 addresses are running out (RIRs allocate IPv4s to regions)



ARIN – American Registry for Internet Numbers








AfriNIC – African Network Information Centre

APNIC – Asia Pacific (most of Asia, Australia, New Zealand) – China in here

LACNIC – Latin America and Caribbean


RIPE NCC – Reseaux IP Europeens (Europe Russian, Middle East and Central Asia)

IPv4 addresses are unfairly allocated

| Country or entity | IP addresses ^[3] | % | Population (mostly 2012) ^[4] | IP addresses per 1000 |
|--|-----------------------------|-------|--|--------------------------|
| <i>World</i> | 4,294,967,296 | 100.0 | 7,021,836,029 | 611.66 |
|  United States | 1,541,605,760 | 35.9 | 313,847,465 | 4,911.96 |
| <i>Bogons</i> | 875,310,464 | 20.4 | | |
|  China | 330,321,408 | 7.7 | 1,343,239,923 | 245.91 |
|  United Kingdom | 123,500,144 | 2.9 | 63,047,162 | 1,958.85 |
|  France | 95,078,032 | 2.2 | 65,630,692 | 1,448.68 |
|  Canada | 79,989,760 | 1.9 | 34,300,083 | 2,332.06 |
|  Italy | 50,999,712 | 1.2 | 61,261,254 | 832.50 |
|  India | 34,685,952 | 0.8 | 1,205,073,612 | 28.78 |

https://en.wikipedia.org/wiki/List_of_countries_by_IPv4_address_allocation

NAT: network address translation



all datagrams *leaving* local network have *same* single source NAT IP address: 138.76.29.7, different source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

NAT: network address translation

motivation: local network uses just one IP address as far as outside world is concerned:


- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

NAT: network address translation

implementation: NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: network address translation



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

NAT: network address translation

- 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
 - routers should only process up to layer 3
 - address shortage should be solved by IPv6
 - violates end-to-end argument (complexity should be at network “ends” not middle)
 - NAT possibility must be taken into account by app designers, e.g., P2P applications

Network Data Plane: outline

4.1 Overview of Network layer

- data plane
- control plane

4.2 What's inside a router

4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing
- network address translation
- IPv6

4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

IPv6: motivation

- *initial motivation*: 32-bit address space soon to be completely allocated.
- additional motivation:
 - header format helps speed processing/forwarding
 - header changes to facilitate QoS

IPv6 datagram format:

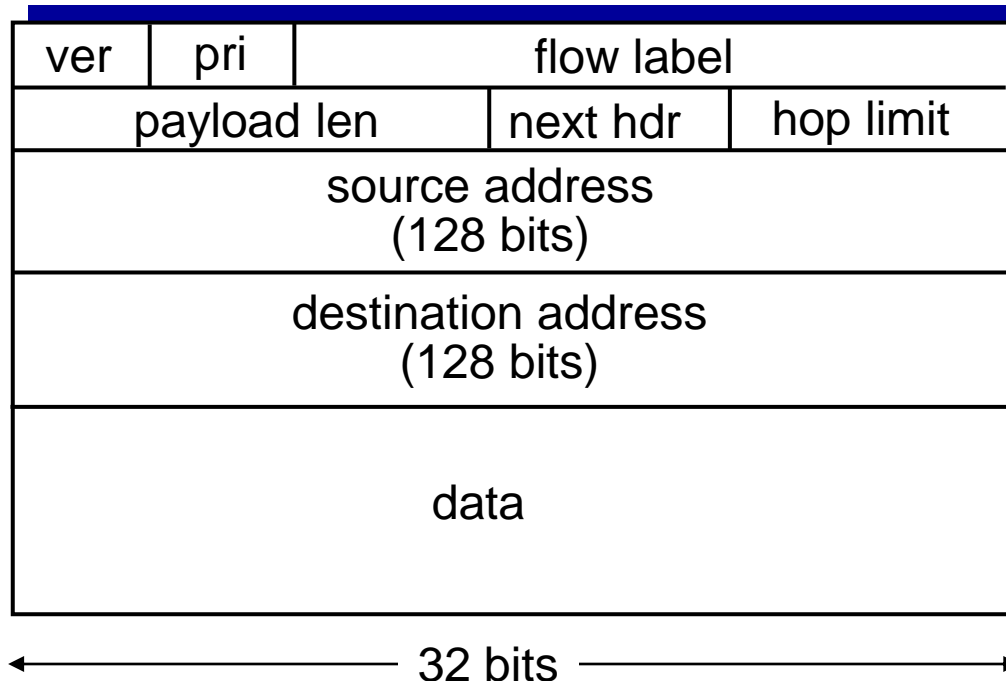
- fixed-length 40 byte header
- no fragmentation allowed

IPv6 datagram format

priority: identify priority among datagrams in flow

flow Label: identify datagrams in same “flow.”

next header: identify upper layer protocol for data



Other changes from IPv4

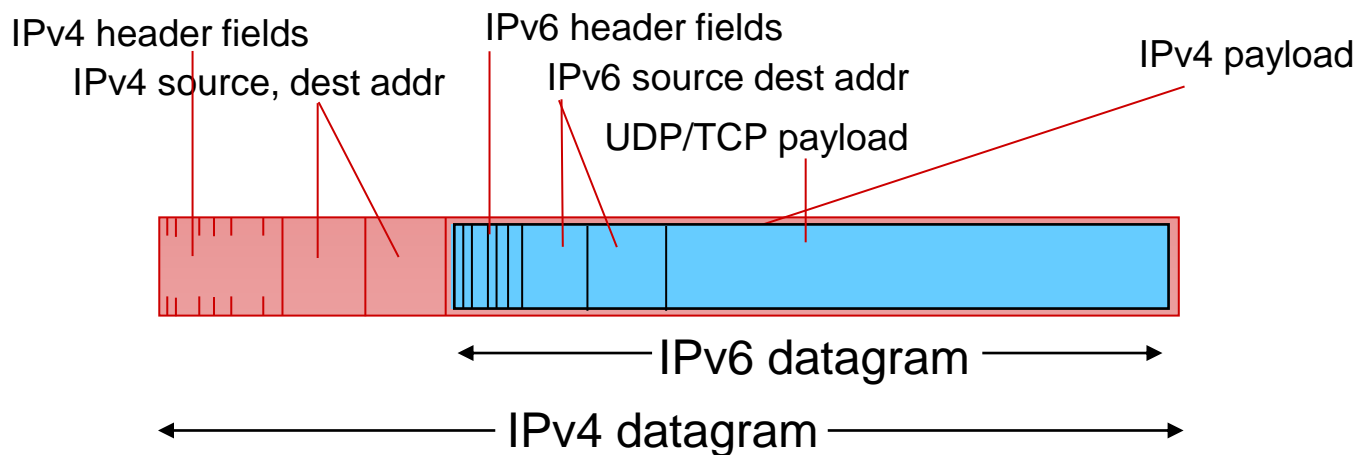
- *checksum*: removed entirely to reduce processing time at each hop
- *options*: allowed, but outside of header, indicated by “Next Header” field
- *ICMPv6*: new version of ICMP (see later lecture)
 - additional message types, e.g. “Packet Too Big”
 - multicast group management functions

IPv4 and IPv6:

Address bits: 32bits and 128bits


Transition from IPv4 to IPv6

- not all routers can be upgraded simultaneously
 - how will network operate with mixed IPv4 and IPv6 routers?
- **tunneling**: IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers

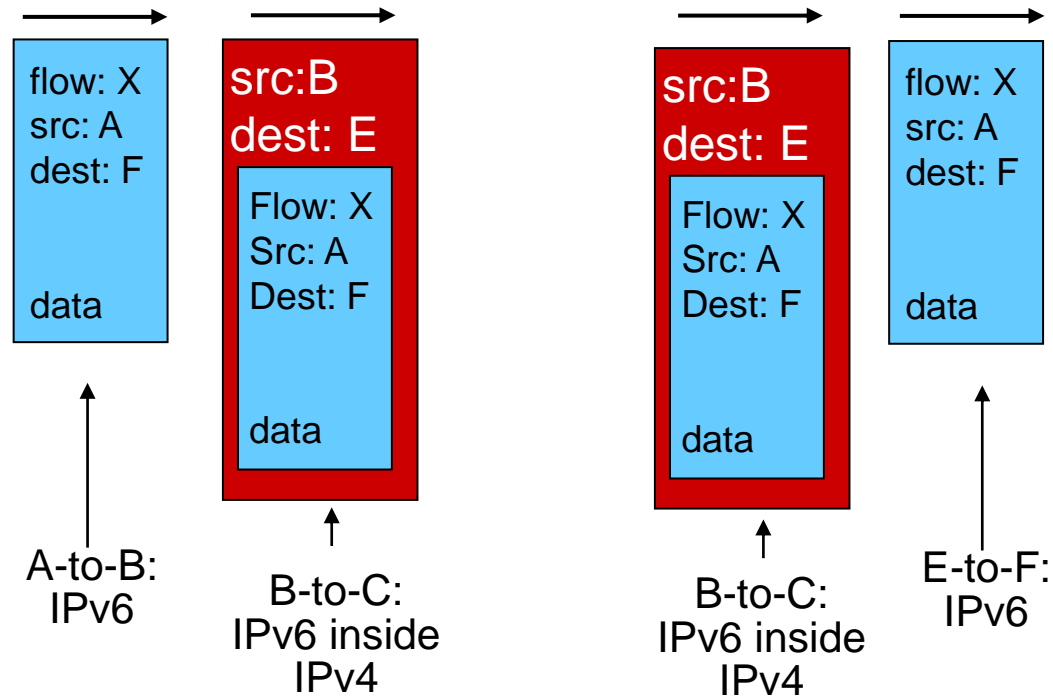



Tunneling

logical view:



physical view:



Test your understanding

- How many bits in an IPv6 address?
 - a) 32 bits
 - b) 64 bits
 - c) 128 bits
- Which of the following not in the IPv6 header?
 - a) Checksum
 - b) Next Header
 - c) Hop Limit

Test your understanding

- How many bits in an IPv6 address?
 - a) 32 bits
 - b) 64 bits
 - c) 128 bits
- Which of the following not in the IPv6 header?
 - a) Checksum
 - b) Next Header
 - c) Hop Limit

Network Data Plane: outline

4.1 Overview of Network layer

- data plane
- control plane

4.2 What's inside a router

4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing
- network address translation
- IPv6

4.4 Generalized Forward and SDN

- match
- action
- OpenFlow examples of match-plus-action in action

What is Software Defined Networking?

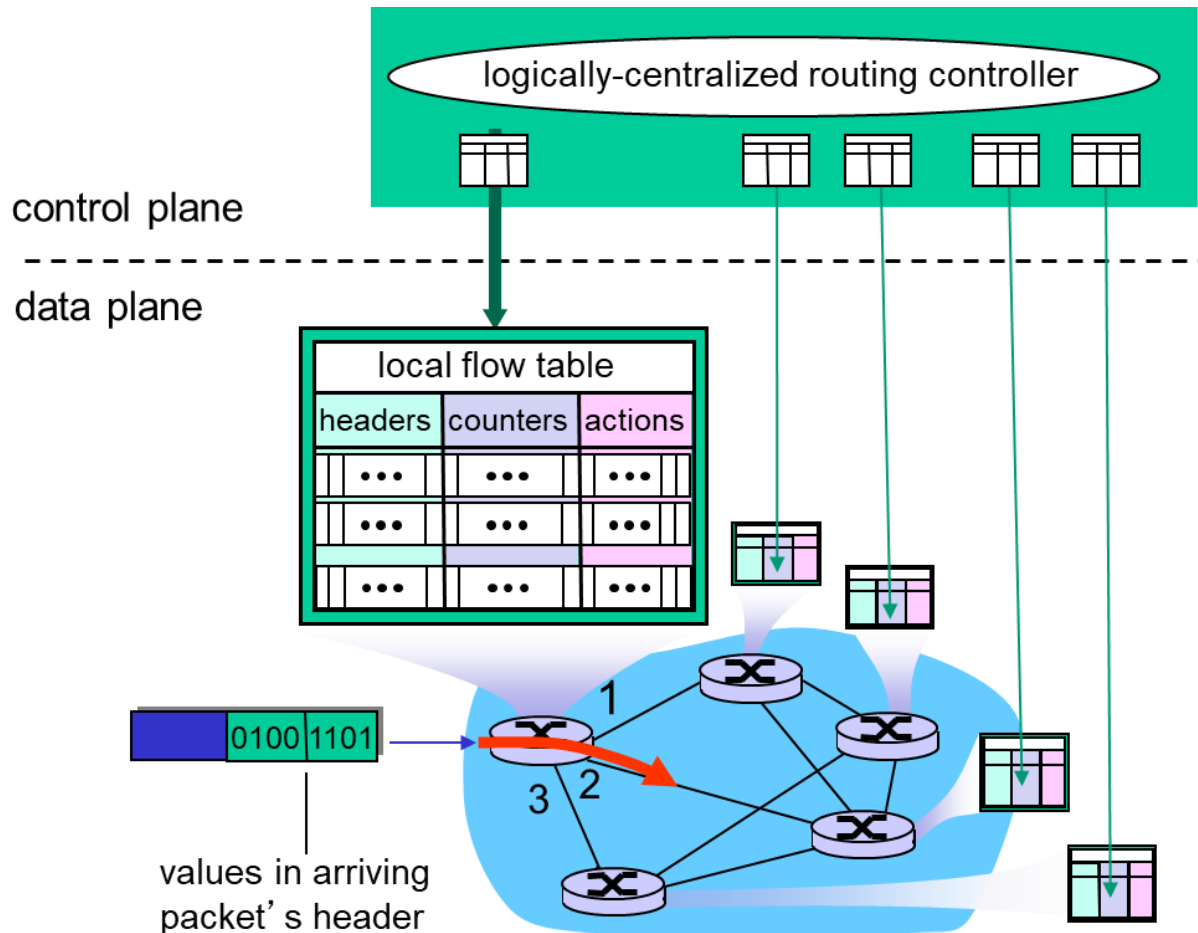
- Traditional routing methods:
 - Data Plane is longest-prefix match forwarding using a forwarding table.
 - The control plane calculates the forwarding table for each router (we will see how later).
 - The forwarding table can forward packets by their IP address and nothing else.
- SDN allows more flexibility in the control and data plane.
 - Program your own control algorithms in language you know (java, python etc).
 - Forwarding can use any part of the packet header.

How does SDN work?

- Data plane:
 - A set of "match-action" rules send by a controller can do many different things to packets. (Forward them, change data etc).
 - Much more flexible (route video packets differently? route private data separately? Drop suspicious data!)
- Control plane (later lectures):
 - Not a distributed system, but a single centralised control point.
 - Programmable, not fixed – you can program the controller in a high-level language that you know.
 - Create your own algorithms and test them on the network without spending a million dollars to create a new hardware router.


Generalized Forwarding and SDN

Each router contains a *flow table* that is computed and distributed by a *logically centralized routing controller*



OpenFlow data plane abstraction

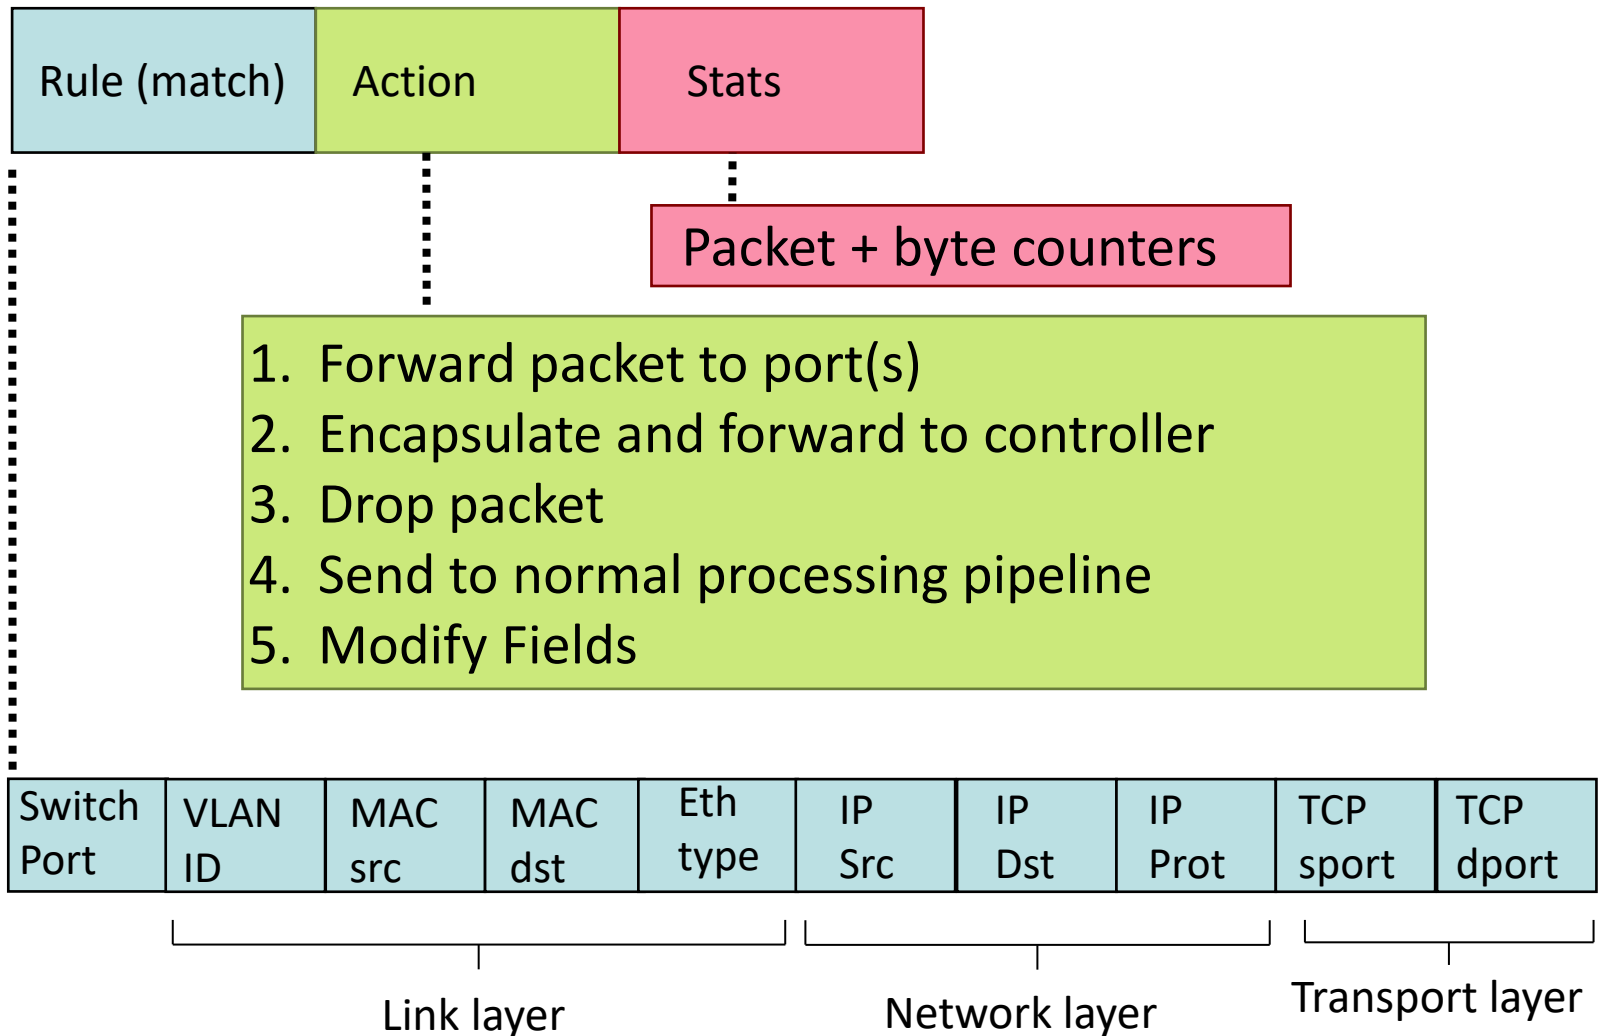
- OpenFlow is a specific SDN protocol
- *flow*: defined by header fields
- generalized forwarding: simple packet-handling rules
 - *Pattern*: match values in packet header fields
 - *Actions: for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
 - *Priority*: disambiguate (tell difference between) overlapping patterns
 - *Counters*: #bytes



* : wildcard

1. src=1.2.*.*, dest=3.4.5.* → drop
2. src = *.*.*.*, dest=3.4.*.* → forward(2)
3. src=10.1.2.3, dest=*.*.*.* → send to controller

OpenFlow: Flow Table Entries



Examples

Destination-based forwarding:

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|--------|----------|---------|-----------|-----------|--------|
| * | * | * | * | * | * | 51.6.0.8 | * | * | * | port6 |

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

Firewall:

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|--------|--------|---------|-----------|-----------|--------|
| * | * | * | * | * | * | * | * | * | 22 | drop |

do not forward (block) all datagrams destined to TCP port 22

| Switch Port | MAC src | MAC dst | Eth type | VLAN ID | IP Src | IP Dst | IP Prot | TCP sport | TCP dport | Action |
|-------------|---------|---------|----------|---------|-------------|--------|---------|-----------|-----------|--------|
| * | * | * | * | * | 128.119.1.1 | * | * | * | * | drop |

do not forward (block) all datagrams sent by host 128.119.1.1

OpenFlow abstraction

- *match+action*: different kinds of devices become one
- Router
 - *match*: longest destination IP prefix
 - *action*: forward out a link
- Switch
 - *match*: destination MAC address
 - *action*: forward or flood
- Firewall
 - *match*: IP addresses and TCP/UDP port numbers
 - *action*: permit or deny
- NAT
 - *match*: IP address and port
 - *action*: rewrite address and port

Network Layer Data Plane: *done!*

4.1 Overview of Network layer: data plane and control plane

4.2 What's inside a router

4.3 IP: Internet Protocol

- datagram format
- fragmentation
- IPv4 addressing
- NAT
- IPv6

4.4 Generalized Forward and SDN

- match plus action
- OpenFlow example

Test your understanding

Which of the following information can be used for “match-action” in OpenFlow?

- a) IP addresses
- b) MAC addresses
- c) Port addresses
- d) Any of the above

Network Control Plane: outline

5.1 introduction

5.2 routing protocols

- link state
- distance vector

5.3 intra-AS routing in the Internet: OSPF

5.4 routing among the ISPs: BGP

5.5 The SDN control plane

5.6 ICMP: The Internet Control Message Protocol

Network-layer functions

Recall: two network-layer functions:

- *forwarding*: move packets from router's input to appropriate router output

data plane

- *routing*: determine route taken by packets from source to destination


control plane

Two approaches to structuring network control plane:

- per-router control (traditional)
- logically centralized control (software defined networking)


Per-router control plane (revision)

Individual routing algorithm components *in each and every router* interact with each other in control plane to compute forwarding tables



Logically centralized control plane (revision)

A distinct (typically remote) controller interacts with local control agents (CAs) in routers to compute forwarding tables



What have we learned?

- NAT (Network Address Translation) is a way of making one IPv4 address usable by many hosts
 - NAT is really common (especially in China)
 - NAT has its problems (getting data back into computer)
- IPv6 improves upon IPv4 in very simple ways.
 - IPv6 deployment slowly growing but not “there” yet.
- Software Defined Networks (for example OpenFlow) is a new technology for forwarding
 - Becoming very popular, extremely powerful
 - We will see later this week how SDN and OpenFlow are used in the control plane.