

# Internet Protocols EBU5403

## Live Lecture B1/B2/B3

Module organiser: Richard Clegg  
([r.clegg@qmul.ac.uk](mailto:r.clegg@qmul.ac.uk))

Michael Chai ([michael.chai@qmul.ac.uk](mailto:michael.chai@qmul.ac.uk))

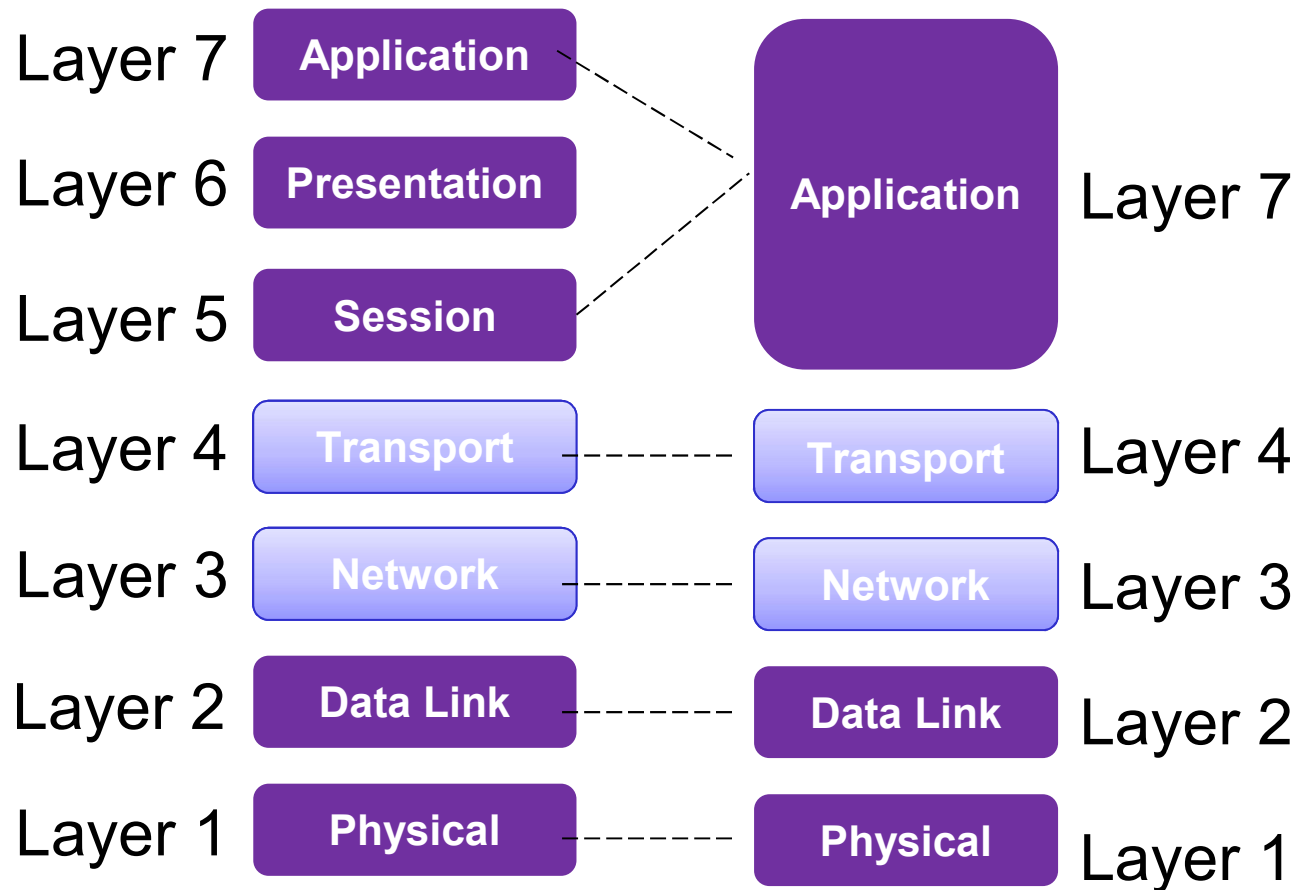
Cunhua Pan([c.pan@qmul.ac.uk](mailto:c.pan@qmul.ac.uk))

	Part 1	Part 2	Part 3	Part 4
Ecommerce + Telecoms 1	Richard Clegg		Cunhua Pan	
Telecoms 2	Michael Chai			

# Structure of course

- Part A
  - Introduction to IP Networks
  - The Transport layer (part I)
- Part B
  - The Transport layer (part II)
  - The Network layer (part I)
  - Class test
- Part C
  - The Network layer (part II)
  - The Data link layer (part I)
  - Router lab tutorial (assessed lab work after this week)
- Part D
  - The Data link layer (part II)
  - Network management and security
  - Class test

# Layers in this life lecture




# Reminder of lecture contents

- Lecture B1
  - Stop and wait vs Go-back-N and Selective Repeat
  - TCP RTT estimate (Jacobson-Karek)
  - TCP fast retransmit
- Lecture B2
  - TCP flow control
  - TCP start up/shut down
  - TCP congestion control
- Lecture B3
  - Data plane and control plane overview
  - Forwarding + longest prefix match


# Link utilisation

- Define link utilisation:
  - A. The time it takes between the sender finishing sending and the sender getting an ACK from the receiver.
  - B. The time it takes to send a packet.
  - C. The proportion of the time that the link is used for sending data.
  - D. The proportion of the time that the link is empty.




# Link utilisation

- Define link utilisation:
  - A. The time it takes between the sender finishing sending and the sender getting an ACK from the receiver. (This is RTT).
  - B. The time it takes to send a packet. (This is transmission delay)
  - C. The proportion of the time that the link is used for sending data.**
  - D. The proportion of the time that the link is empty. (This is  $1 - \text{utilisation}$ ).



# Link utilisation

- Consider data transmission in rdt 3.0. Packets of length 1500 bits are sent over a link which is 100Mb/s. The delay is 1ms.
- What is the utilisation (approx):
  - A. 0.00744
  - B. 0.00515
  - C. 0.00015
  - D. 1.5



# Link utilisation

- Consider data transmission in rdt 3.0. Packets of length 1500 bits are sent over a link which is 100Mb/s. The delay is 1ms.

- What is the utilisation (approx):

**A. 0.00744**

B. 0.00515

C. 0.00015


D. 1.5

$L = 1500 \text{ bits}$   $R = 100,000,000 \text{ bits/sec}$

$L/R = 0.000015 \text{ seconds}$

$RTT = 0.002 \text{ seconds}$

$U = (L/R) / (L/R + RTT) = 0.000015 / 0.002015 = 0.00744..$






# Pipelining to increase utilisation

Instead of sending packets one after another packets are sent in batches of  $N$ . The link is not full and no losses are experienced. What is effect on utilisation


- A. Increases by a factor of  $N$
- B. Stays the same
- C. Becomes three times as much
- D. Decreases to  $1/N$




# Pipelining to increase utilisation

Instead of sending packets one after another packets are sent in batches of  $N$ . The link is not full and no losses are experienced. What is effect on utilisation

- A. Increases by a factor of  $N$**
- B. Stays the same
- C. Becomes three times as much
- D. Decreases to  $1/N$




# rdt3.0: stop-and-wait operation



$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$

# Pipelining: increased utilization



# Pipelined protocols: overview

## Go-back-N:

- sender can have up to N unacked packets in pipeline
- receiver only sends *cumulative ACK*
  - doesn't ACK packet if there's a gap
- sender has timer for oldest unacked packet
  - when timer expires, retransmit *all* unacked packets

## Selective Repeat:

- sender can have up to N packets which it has not seen ACK for in “pipeline”
- Receiver sends *individual ACK* for each packet
- sender maintains timer for each packet with no ACK
  - when timer expires, retransmit only that unacked packet

# Go back N

- Which of these statements is true for go-back-N receiver if it gets packets 1,2,3,5
  - A) Should send ACKs 1,2,3,5
  - B) Should send packets 1,2,3,5 to the application
  - C) Should send no packets to application because of missing data
  - D) Should discard packet 5.

# Go back N

- Which of these statements is true for go-back-N receiver if it gets packets 1,2,3,5

A) Should send ACKs 1,2,3,5

No ACK if there is a gap

B) Should send packets 1,2,3,5 to the application

Packet 5 would be out of order so don't send.

C) Should send no packets to application because of missing data

Apart from anything else it does not know 4 is missing when it receives 1,2 and 3.

**D) Should discard packet 5.**

# Selective repeat

- Which is true for selective repeat receiver if it gets packets 1,2,3,5
  - A) Should send ACKs 1,2,3,5
  - B) Should send packets 1,2,3,5 to the application
  - C) Should send no packets to application because of missing data
  - D) Should discard packet 5.



# Selective repeat

- Which of these statements is true for go-back-N receiver if it gets packets 1,2,3,5

**A) Should send ACKs 1,2,3,5**

Yes – it acknowledges every packet

**B) Should send packets 1,2,3,5 to the application**

Packet 5 would be out of order so don't send.

**C) Should send no packets to application because of missing data**

Apart from anything else it does not know 4 is missing when it receives 1,2 and 3.

**D) Should discard packet 5.**

Keeps it in buffer and waits for 4.

# Advantages each scheme

- Go Back N
  - Simple to implement
  - Only a single timer
  - One ACK can acknowledge many packets
- Selective repeat
  - Minimises resending of packets (better use of network)
  - Individual ACKs give more information to sender

# Jacobsen-Karel Algorithm

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- $\alpha = 0.125$  (1/8)
- The previous estimate of RTT was 80ms.
- A new sample arrives with RTT = 160ms.
- What is the new estimated RTT?

# Jacobsen-Karel Algorithm

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- $\alpha = 0.125$  (1/8)
- The previous estimate of RTT was 80ms.
- A new sample arrives with RTT = 160ms.
- What is the new estimated RTT?
  
- $(1 - \alpha) = 7/8$
- $\text{EstimatedRTT} = 7/8 * 80 + 1/8 * 160\text{ms}$
- $\text{EstimatedRTT} = 90\text{ms}$

# TCP flow control

Which statements about flow control are true:

- A) The receive window (rwnd) is sent over the network to the sender.
- B) The receive window measures how much congestion is on the network
- C) Flow control reduces network congestion
- D) Sender sees rwnd=5000bytes. It has 4 UNACK'd packets of 1500bytes. It can send 1 more.

# TCP flow control

Which statements about flow control are true:

- A) The receive window (rwnd) is sent over the network to the sender.

Yes – it is in the TCP header.

- B) The receive window measures how much congestion is on the network

No, it measures the receive buffer

- C) Flow control reduces network congestion

Not deliberately, cannot see network congestion

- D) Sender sees rwnd=5000bytes. It has 4 UNACK'd packets of 1500bytes. It can send 1 more.

No – 6000 bytes already send

---




100

**flow control**  
receiver controls sender, so  
sender will not overflow  
receiver buffer by transmitting  
too much, too fast



# TCP flow control

- receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- guarantees receive buffer will not overflow



**rwnd** = receive window



# TCP handshaking

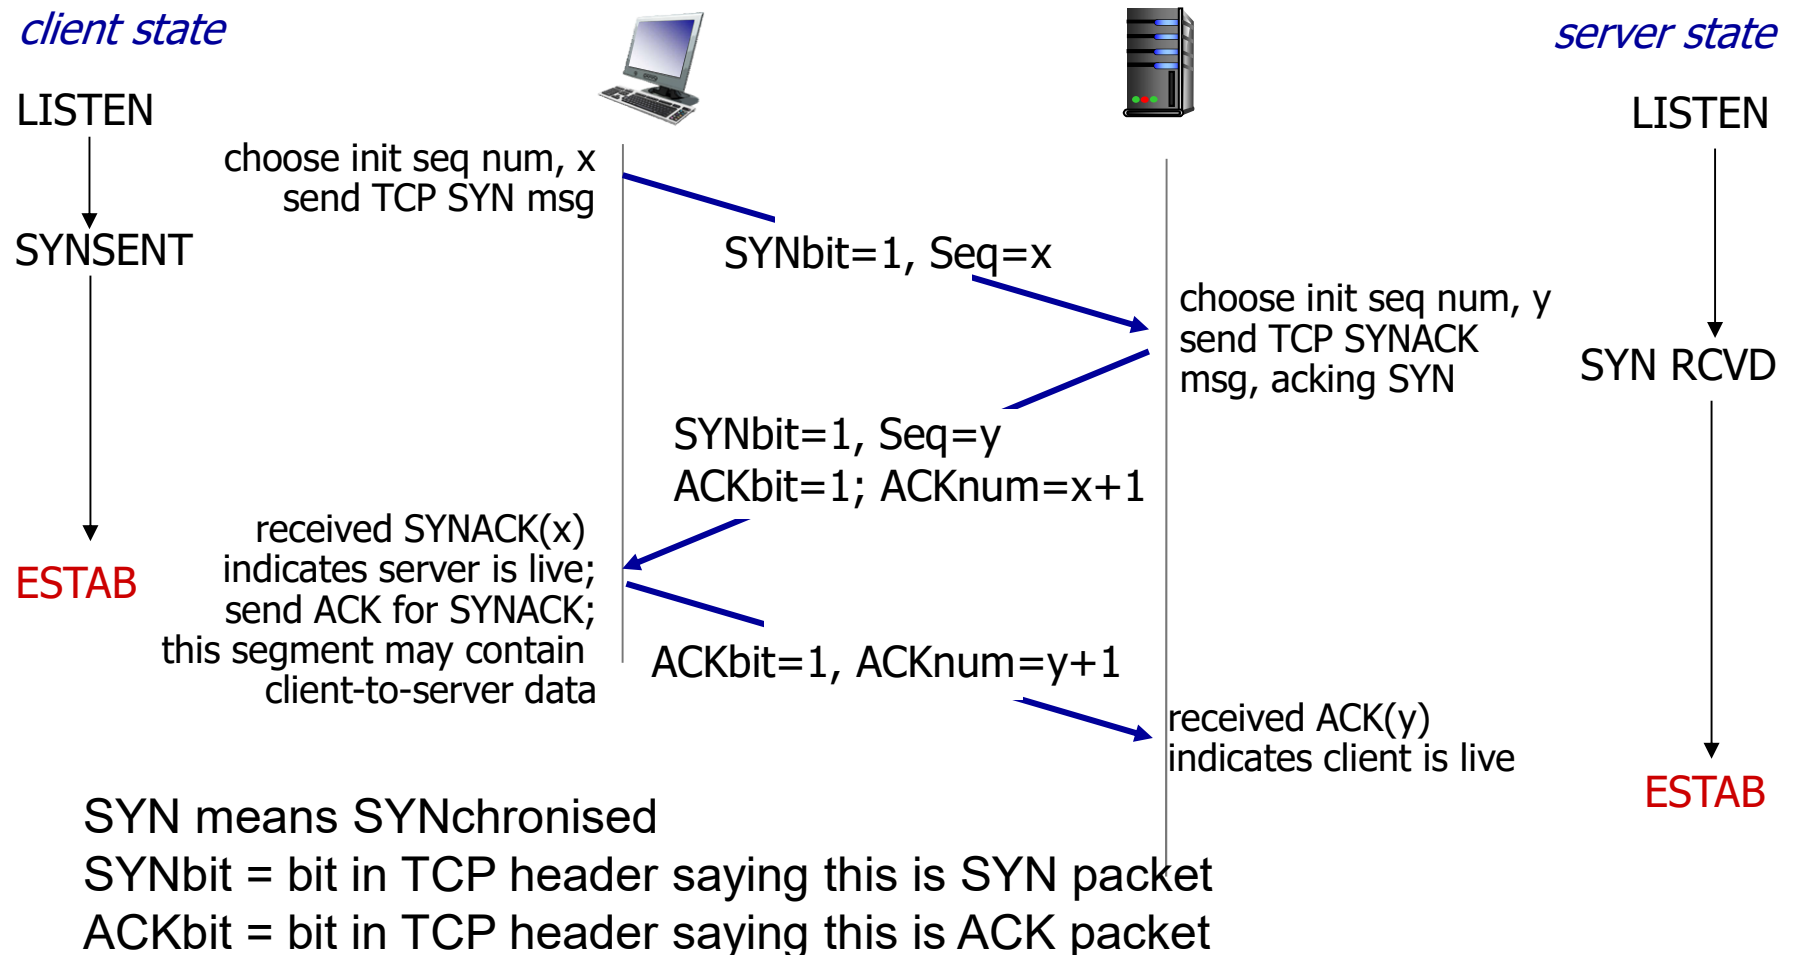
- What is the sequence for the TCP three way handshake to open a connection?
- To close a connection the sequence is FIN, FIN-ACK, ACK. Why might the FIN-ACK be split into two packets?

# TCP handshaking

- What is the sequence for the TCP three way handshake to open a connection?
  - SYN SYN-ACK ACK
- To close a connection the sequence is FIN, FIN-ACK, ACK. Why might the FIN-ACK be split into two packets?

The closing computer may have data left to send.

# TCP 3-way handshake



# TCP: closing a connection

*client state*

ESTAB

FIN\_WAIT\_1

FIN\_WAIT\_2

TIMED\_WAIT


CLOSED

`clientSocket.close()`

can no longer  
send but can  
receive data

wait for server  
close

timed wait  
for  $2 * \text{max}$   
segment lifetime



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still  
send data

can no longer  
send data

*server state*

ESTAB

CLOSE\_WAIT

LAST\_ACK

CLOSED

# TCP congestion control

Which statements about congestion control are true:

- A) The congestion window (cwnd) is sent over the network to the sender.
- B) AIMD means that cwnd increases only slowly when there is no loss
- C) Congestion control has problems because of loss
- D) Congestion control reduces network congestion

# TCP congestion control

Which statements about congestion control are true:

- A) The congestion window (cwnd) is sent over the network to the sender.

No, it is calculated by the sender.

- B) AIMD means that cwnd increases only slowly when there is no loss

Yes – it increases slowly decreases quickly.

- C) Congestion control has problems because of loss

Loss is fundamental to most congestion control

- D) Congestion control reduces network congestion

Yes – this is its very aim.

# TCP mechanisms

Name these three TCP mechanisms

- When I see four Acknowledgement packets with the same number I retransmit a packet.
- When cwnd is small I double cwnd every RTT where no loss is experienced.
- When I think there is no congestion I increase the window size slowly, when I think there is congestion I decrease the window size rapidly


# TCP mechanisms

Name these three TCP mechanisms

- When I see four Acknowledgement packets with the same number I retransmit a packet.
- Fast retransmit
- When cwnd is small I double cwnd every RTT where no loss is experienced.
- Slow start
- When I think there is no congestion I increase the window size slowly, when I think there is congestion I decrease the window size rapidly
- AIMD (Additive Increase Multiplicative Decrease)

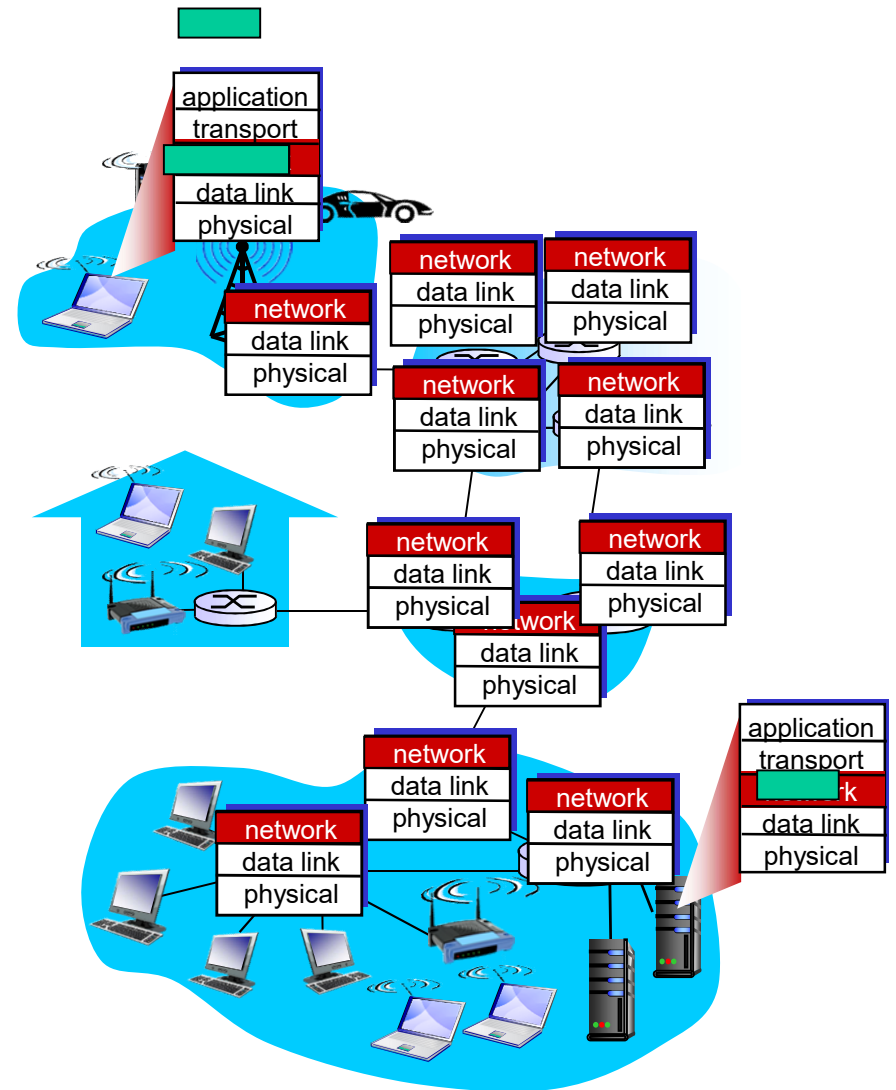


# TCP: switching from slow start to Congestion Avoidance



# Network layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in *every* host, router
- router examines header fields in all IP datagrams passing through it




# Network layer: data plane, control plane

## *Data plane*

- local, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
- forwarding function

values in arriving  
packet header



## *Control plane*

- network-wide logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
  - *traditional routing algorithms*: implemented in routers
  - *software-defined networking (SDN)*: implemented in (remote) servers

# Data plane or control plane

- Finding the correct route for a packet
- Forwarding a packet using Longest Prefix Match
- An SDN controller
- A router

# Data plane or control plane

- Finding the correct route for a packet
  - Control plane
- Forwarding a packet using Longest Prefix Match
  - Data plane – takes place locally using forwarding table
- An SDN controller
  - Control plane (it is in the name)
- A traditional router
  - Both (a router has forwarding and routing)

# Longest prefix matching which interface do the packets go to

Destination Address Range	Link
11001000 00010111 01111*** *****	A
11001000 00010111 01111000 *****	B
11001000 00010111 00011*** *****	C
otherwise	D

Dest: 11001000 00010111 01111000 10101010

Dest: 11001000 00010111 01111100 10101010

Dest: 11001000 00010111 00111000 10101010

# Longest prefix matching which interface do the packets go to

Destination Address Range	Link
11001000 00010111 01111*** *****	A
11001000 00010111 01111000 *****	B
11001000 00010111 00011*** *****	C
otherwise	D

Dest: 11001000 00010111 01111000 10101010

B

Dest: 11001000 00010111 01111100 10101010

A

Dest: 11001000 00010111 00111000 10101010

D

# Longest prefix matching

- we'll see *why* longest prefix matching is used shortly, when we study addressing
- longest prefix matching: often performed using ternary content addressable memories (TCAMs) specialised very high speed memory
  - *content addressable*: present address to TCAM: retrieve address in one clock cycle, regardless of table size
  - Cisco Catalyst: can up ~1M routing table entries in TCAM



# What have we learned

- Throughput and pipelined protocols
  - Stop and wait vs Go-back-N and Selective Repeat
- TCP mechanisms
  - TCP RTT estimate (Jacobson-Karek)
  - TCP fast retransmit
  - TCP flow control
  - TCP start up/shut down
  - TCP congestion control
- Network layer
  - Data plane and control plane overview
  - Forwarding + longest prefix match